

## 7) Alignment & Reproducibility Report (對齊與可重現性報告)

Date: 2025-12-24

### Objective

Ensure that training, distillation, and evaluation are reproducible and that data/label alignment is consistent end-to-end.

### Alignment artifacts (data)

- Unified manifest validation:
  - `outputs/manifest\_validation\_all\_with\_expw.json` (466,284 rows, 0 missing paths, 0 bad labels)
  - `outputs/manifest\_validation.json` (374,491 rows, 0 missing paths)
- Cleaning provenance:
  - `Training\_data\_cleaned/clean\_report.json`

### Alignment artifacts (distillation)

- Softlabel export artifacts are explicit and indexable:
  - `outputs/softlabels/.../softlabels.npz`
  - `outputs/softlabels/.../softlabels\_index.jsonl`
- Teacher vs student decoupling:
  - Student training consumes saved teacher logits (no need to run teacher ONNX at student training time).

### Checkpoint / run provenance

- Teacher + student runs save reproducible artifacts:
  - `best.pt` checkpoints
  - `history.json` curves
  - `reliabilitymetrics.json` and `calibration.json` evaluations
- Provenance helpers:

- `scripts/inspect\_checkpoint\_provenance.py`
- `\_ckpt\_provenance\_\*.txt` snapshots

## Windows reproducibility notes

- Windows multiprocessing can be fragile (DataLoader, shared mappings). The project uses safer defaults and PowerShell orchestration for stable runs.

## DKD resume incident (and fix)

- Root cause: DKD resumed from KD at a `start\_epoch` greater than the configured total `--epochs`, resulting in a no-op training loop and an empty output.
- Fix: compute DKD total epochs as (resume\_epoch + 1 + extra\_dkd\_epochs) inside the runner.
- This was documented in the weekly log and validated by a successful DKD-only rerun producing full artifacts.

## Next steps

- Add a single “run manifest” per experiment (args + git commit hash + environment packages) to make reproducing any run one command.